

# Deep Q-Learning

Install dependencies for AI gym to run properly (shouldn't take more than a minute). If running on google cloud or running locally, only need to run once. Colab may require installing everytime the vm shuts down.

```
In [1]: # !pip3 install gym pyvirtualdisplay
# !sudo apt-get install -y xvfb python-opengl ffmpeg
```

```
In [2]: # !pip3 install --upgrade setuptools --user
# !pip3 install ez_setup
# !pip3 install gym[atari]
```

For this assignment we will implement the Deep Q-Learning algorithm with Experience Replay as described in breakthrough paper "**Playing Atari with Deep Reinforcement Learning**". We will train an agent to play the famous game of **Breakout**.

```
In [1]: %matplotlib inline

import sys
import gym
import torch
import pylab
import random
import numpy as np
from collections import deque
from datetime import datetime
from copy import deepcopy
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.autograd import Variable
from utils import find_max_lives, check_live, get_frame, get_init_state
from model import DQN
from config import *

import matplotlib.pyplot as plt
# %load_ext autoreload
# %autoreload 2
```

```
In [4]: # !pip3 install -q gym[atari]
# !pip install -q autorom[accept-rom-license]
```

```
In [5]: env = gym.make('ALE/Breakout-v5')
state = env.reset()
#print(state)
```

```
A.L.E: Arcade Learning Environment (version 0.8.1+53f58b7)
[Powered by Stella]
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/core.py:317: DeprecationWarning: WARN: Initializing wrapper in old step API which returns one bool instead of two. It is recommended to set `new_step_api=True` to use new step API. This will be the default behaviour in future.
  deprecation(
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/wrappers/step_api_compatibility.py:39: DeprecationWarning: WARN: Initializing environment in old step API which returns one bool instead of two. It is recommended to set `new_step_api=True` to use new step API. This will be the default behaviour in future.
  deprecation(
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/utils/passive_env_checker.py:190: UserWarning: WARN: Future gym versions will require that `Env.reset` can be passed `return_info` to return information from the environment resetting.
  logger.warn(
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/utils/passive_env_checker.py:137: UserWarning: WARN: The obs returned by the `reset()` method was expecting a numpy array, actual type: <class 'tuple'>
  logger.warn(
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/spaces/box.py:226: UserWarning: WARN: Casting input x to numpy array.
  logger.warn("Casting input x to numpy array.")
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/utils/passive_env_checker.py:167: UserWarning: WARN: The obs returned by the `reset()` method is not within the observation space with exception: setting an array element with a sequence. The requested array has an inhomogeneous shape after 1 dimensions. The detected shape was (2,) + inhomogeneous part.
  logger.warn(f"{pre} is not within the observation space with exception: {e}")
```

## Understanding the environment

In the following cell, we initialize our game of **Breakout** and you can see how the environment looks like. For further documentation of the of the environment refer to <https://gym.openai.com/envs>.

In breakout, we will use 3 actions "fire", "left", and "right". "fire" is only used to reset the game when a life is lost, "left" moves the agent left and "right" moves the agent right.

```
In [6]: # env = gym.make('BreakoutDeterministic-v4')
# state = env.reset()
```

```
In [7]: number_lives = find_max_lives(env)
state_size = env.observation_space.shape
action_size = 3 #fire, left, and right
```

```
/home/aalkhami/miniconda3/envs/myenv/lib/python3.9/site-packages/gym/utils/passive_env_checker.py:241: DeprecationWarning: `np.bool8` is a deprecated alias for `np.bool_`. (Deprecated NumPy 1.24)
    if not isinstance(terminated, (bool, np.bool8)):
```

## Creating a DQN Agent

Here we create a DQN Agent. This agent is defined in the **agent.py**. The corresponding neural network is defined in the **model.py**. Once you've created a working DQN agent, use the code in agent.py to create a double DQN agent in **agent\_double.py**. Set the flag "double\_dqn" to True to train the double DQN agent.

**Evaluation Reward** : The average reward received in the past 100 episodes/games.

**Frame** : Number of frames processed in total.

**Memory Size** : The current size of the replay memory.

```
In [8]: double_dqn = False # set to True if using double DQN agent

if double_dqn:
    from agent_double import Agent
else:
    from agent import Agent

agent = Agent(action_size)
evaluation_reward = deque(maxlen=evaluation_reward_length)
frame = 0
memory_size = 0
```

## Main Training Loop

In this training loop, we do not render the screen because it slows down training significantly. To watch the agent play the game, run the code in next section "Visualize Agent Performance"

```
In [9]: rewards, episodes = [], []
best_eval_reward = 0
for e in range(EPISODES):
    done = False
    score = 0
```

```

history = np.zeros([5, 84, 84], dtype=np.uint8)
step = 0
d = False
state = env.reset()
next_state = state
life = number_lives

get_init_state(history, state)

while not done:
    step += 1
    frame += 1

    # Perform a fire action if ball is no longer on screen to continue c
    if step > 1 and len(np.unique(next_state[:189] == state[:189])) < 2:
        action = 0
    else:
        action = agent.get_action(np.float32(history[:4, :, :]) / 255.)
    state = next_state
    next_state, reward, done, info = env.step(action + 1)

    frame_next_state = get_frame(next_state)
    history[4, :, :] = frame_next_state
    terminal_state = check_live(life, info['lives'])

    life = info['lives']
    r = np.clip(reward, -1, 1)
    r = reward

    # Store the transition in memory
    agent.memory.push(deepcopy(frame_next_state), action, r, terminal_st
    # Start training after random sample generation
    if (frame >= train_frame):
        agent.train_policy_net(frame)
        # Update the target network only for Double DQN only
        if double_dqn and (frame % update_target_network_frequency) == 0:
            agent.update_target_net()
    score += reward
    history[:4, :, :] = history[1:, :, :]

    if done:
        evaluation_reward.append(score)
        rewards.append(np.mean(evaluation_reward))
        episodes.append(e)
        pylab.plot(episodes, rewards, 'b')
        pylab.xlabel('Episodes')
        pylab.ylabel('Rewards')
        pylab.title('Episodes vs Reward')
        pylab.savefig("./save_graph/breakout_dqn.png") # save graph for

    # every episode, plot the play time
    print("episode:", e, " score:", score, " memory length:",
          len(agent.memory), " epsilon:", agent.epsilon, " steps:

```

```

        "lr:", agent.optimizer.param_groups[0]['lr'], "eval

# if the mean of scores of last 100 episode is bigger than 8 save
### Change this save condition to whatever you prefer ###
if np.mean(evaluation_reward) >= 8 and np.mean(evaluation_reward
    torch.save(agent.policy_net, "./save_model/breakout_dqn.pth"
    best_eval_reward = np.mean(evaluation_reward)

```

```

/tmp/ipykernel_975387/1450953243.py:21: DeprecationWarning: elementwise comp
arison failed; this will raise an error in the future.

```

```

    if step > 1 and len(np.unique(next_state[:189] == state[:189])) < 2:
episode: 0    score: 2.0    memory length: 198    epsilon: 1.0    steps: 198
lr: 0.0001    evaluation reward: 2.0
episode: 1    score: 3.0    memory length: 445    epsilon: 1.0    steps: 247
lr: 0.0001    evaluation reward: 2.5
episode: 2    score: 0.0    memory length: 568    epsilon: 1.0    steps: 123
lr: 0.0001    evaluation reward: 1.6666666666666667
episode: 3    score: 3.0    memory length: 815    epsilon: 1.0    steps: 247
lr: 0.0001    evaluation reward: 2.0
episode: 4    score: 2.0    memory length: 1012    epsilon: 1.0    steps: 197
lr: 0.0001    evaluation reward: 2.0
episode: 5    score: 2.0    memory length: 1210    epsilon: 1.0    steps: 198
lr: 0.0001    evaluation reward: 2.0
episode: 6    score: 3.0    memory length: 1438    epsilon: 1.0    steps: 228
lr: 0.0001    evaluation reward: 2.142857142857143
episode: 7    score: 1.0    memory length: 1606    epsilon: 1.0    steps: 168
lr: 0.0001    evaluation reward: 2.0
episode: 8    score: 5.0    memory length: 1913    epsilon: 1.0    steps: 307
lr: 0.0001    evaluation reward: 2.3333333333333335
episode: 9    score: 2.0    memory length: 2111    epsilon: 1.0    steps: 198
lr: 0.0001    evaluation reward: 2.3
episode: 10   score: 3.0    memory length: 2339    epsilon: 1.0    steps: 228
lr: 0.0001    evaluation reward: 2.3636363636363638
episode: 11   score: 1.0    memory length: 2511    epsilon: 1.0    steps: 172
lr: 0.0001    evaluation reward: 2.25
episode: 12   score: 3.0    memory length: 2757    epsilon: 1.0    steps: 246
lr: 0.0001    evaluation reward: 2.3076923076923075
episode: 13   score: 4.0    memory length: 3016    epsilon: 1.0    steps: 259
lr: 0.0001    evaluation reward: 2.4285714285714284
episode: 14   score: 1.0    memory length: 3188    epsilon: 1.0    steps: 172
lr: 0.0001    evaluation reward: 2.3333333333333335
episode: 15   score: 0.0    memory length: 3311    epsilon: 1.0    steps: 123
lr: 0.0001    evaluation reward: 2.1875
episode: 16   score: 1.0    memory length: 3480    epsilon: 1.0    steps: 169
lr: 0.0001    evaluation reward: 2.1176470588235294
episode: 17   score: 1.0    memory length: 3649    epsilon: 1.0    steps: 169
lr: 0.0001    evaluation reward: 2.0555555555555554
episode: 18   score: 1.0    memory length: 3800    epsilon: 1.0    steps: 151
lr: 0.0001    evaluation reward: 2.0
episode: 19   score: 2.0    memory length: 4018    epsilon: 1.0    steps: 218
lr: 0.0001    evaluation reward: 2.0
episode: 20   score: 3.0    memory length: 4265    epsilon: 1.0    steps: 247
lr: 0.0001    evaluation reward: 2.0476190476190474

```

episode: 21	score: 2.0	memory length: 4452	epsilon: 1.0	steps: 187
lr: 0.0001	evaluation reward: 2.0454545454545454			
episode: 22	score: 0.0	memory length: 4574	epsilon: 1.0	steps: 122
lr: 0.0001	evaluation reward: 1.9565217391304348			
episode: 23	score: 2.0	memory length: 4792	epsilon: 1.0	steps: 218
lr: 0.0001	evaluation reward: 1.9583333333333333			
episode: 24	score: 0.0	memory length: 4914	epsilon: 1.0	steps: 122
lr: 0.0001	evaluation reward: 1.88			
episode: 25	score: 1.0	memory length: 5085	epsilon: 1.0	steps: 171
lr: 0.0001	evaluation reward: 1.8461538461538463			
episode: 26	score: 0.0	memory length: 5208	epsilon: 1.0	steps: 123
lr: 0.0001	evaluation reward: 1.7777777777777777			
episode: 27	score: 1.0	memory length: 5377	epsilon: 1.0	steps: 169
lr: 0.0001	evaluation reward: 1.75			
episode: 28	score: 3.0	memory length: 5602	epsilon: 1.0	steps: 225
lr: 0.0001	evaluation reward: 1.793103448275862			
episode: 29	score: 2.0	memory length: 5821	epsilon: 1.0	steps: 219
lr: 0.0001	evaluation reward: 1.8			
episode: 30	score: 0.0	memory length: 5943	epsilon: 1.0	steps: 122
lr: 0.0001	evaluation reward: 1.7419354838709677			
episode: 31	score: 4.0	memory length: 6257	epsilon: 1.0	steps: 314
lr: 0.0001	evaluation reward: 1.8125			
episode: 32	score: 0.0	memory length: 6379	epsilon: 1.0	steps: 122
lr: 0.0001	evaluation reward: 1.7575757575757576			
episode: 33	score: 0.0	memory length: 6502	epsilon: 1.0	steps: 123
lr: 0.0001	evaluation reward: 1.7058823529411764			
episode: 34	score: 1.0	memory length: 6673	epsilon: 1.0	steps: 171
lr: 0.0001	evaluation reward: 1.6857142857142857			
episode: 35	score: 2.0	memory length: 6873	epsilon: 1.0	steps: 200
lr: 0.0001	evaluation reward: 1.6944444444444444			
episode: 36	score: 0.0	memory length: 6996	epsilon: 1.0	steps: 123
lr: 0.0001	evaluation reward: 1.6486486486486487			
episode: 37	score: 0.0	memory length: 7119	epsilon: 1.0	steps: 123
lr: 0.0001	evaluation reward: 1.605263157894737			
episode: 38	score: 1.0	memory length: 7270	epsilon: 1.0	steps: 151
lr: 0.0001	evaluation reward: 1.5897435897435896			
episode: 39	score: 3.0	memory length: 7539	epsilon: 1.0	steps: 269
lr: 0.0001	evaluation reward: 1.625			
episode: 40	score: 2.0	memory length: 7737	epsilon: 1.0	steps: 198
lr: 0.0001	evaluation reward: 1.6341463414634145			
episode: 41	score: 2.0	memory length: 7935	epsilon: 1.0	steps: 198
lr: 0.0001	evaluation reward: 1.6428571428571428			
episode: 42	score: 0.0	memory length: 8058	epsilon: 1.0	steps: 123
lr: 0.0001	evaluation reward: 1.6046511627906976			
episode: 43	score: 1.0	memory length: 8227	epsilon: 1.0	steps: 169
lr: 0.0001	evaluation reward: 1.5909090909090908			
episode: 44	score: 0.0	memory length: 8349	epsilon: 1.0	steps: 122
lr: 0.0001	evaluation reward: 1.5555555555555556			
episode: 45	score: 1.0	memory length: 8517	epsilon: 1.0	steps: 168
lr: 0.0001	evaluation reward: 1.5434782608695652			
episode: 46	score: 2.0	memory length: 8720	epsilon: 1.0	steps: 203
lr: 0.0001	evaluation reward: 1.553191489361702			
episode: 47	score: 3.0	memory length: 8968	epsilon: 1.0	steps: 248


episode: 3435	score: 10.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 466	lr: 2.62144000000000017e-08	evaluation reward: 8.64
episode: 3436	score: 14.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 597	lr: 2.62144000000000017e-08	evaluation reward: 8.69
episode: 3437	score: 8.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 424	lr: 2.62144000000000017e-08	evaluation reward: 8.62
episode: 3438	score: 6.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 359	lr: 2.62144000000000017e-08	evaluation reward: 8.55
episode: 3439	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 497	lr: 2.62144000000000017e-08	evaluation reward: 8.54
episode: 3440	score: 11.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 502	lr: 2.62144000000000017e-08	evaluation reward: 8.58
episode: 3441	score: 6.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 340	lr: 2.62144000000000017e-08	evaluation reward: 8.54
episode: 3442	score: 13.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 633	lr: 2.62144000000000017e-08	evaluation reward: 8.61
episode: 3443	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 454	lr: 2.62144000000000017e-08	evaluation reward: 8.63
episode: 3444	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 507	lr: 2.62144000000000017e-08	evaluation reward: 8.64
episode: 3445	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 477	lr: 2.62144000000000017e-08	evaluation reward: 8.7
episode: 3446	score: 16.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 685	lr: 2.62144000000000017e-08	evaluation reward: 8.8
episode: 3447	score: 8.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 434	lr: 2.62144000000000017e-08	evaluation reward: 8.78
episode: 3448	score: 8.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 386	lr: 2.62144000000000017e-08	evaluation reward: 8.8
episode: 3449	score: 12.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 563	lr: 2.62144000000000017e-08	evaluation reward: 8.84
episode: 3450	score: 12.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 630	lr: 2.62144000000000017e-08	evaluation reward: 8.76
episode: 3451	score: 14.0	memory length: 1000000	epsilon: 0.0099980200
08555413	steps: 588	lr: 2.62144000000000017e-08	evaluation reward: 8.83
episode: 3452	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 425	lr: 2.62144000000000017e-08	evaluation reward: 8

```
.86
episode: 3453  score: 9.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 455  lr: 2.62144000000000017e-08  evaluation reward: 8
.88
episode: 3454  score: 14.0  memory length: 1000000  epsilon: 0.0099980200
08555413  steps: 562  lr: 2.62144000000000017e-08  evaluation reward:
8.91
episode: 3455  score: 18.0  memory length: 1000000  epsilon: 0.0099980200
08555413  steps: 678  lr: 2.62144000000000017e-08  evaluation reward:
8.99
episode: 3456  score: 8.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 407  lr: 2.62144000000000017e-08  evaluation reward: 8
.97
episode: 3457  score: 12.0  memory length: 1000000  epsilon: 0.0099980200
08555413  steps: 435  lr: 2.62144000000000017e-08  evaluation reward:
8.94
episode: 3458  score: 16.0  memory length: 1000000  epsilon: 0.0099980200
08555413  steps: 625  lr: 2.62144000000000017e-08  evaluation reward:
9.0
episode: 3459  score: 6.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 345  lr: 2.62144000000000017e-08  evaluation reward: 8
.94
episode: 3460  score: 9.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 451  lr: 2.62144000000000017e-08  evaluation reward: 8
.92
episode: 3461  score: 9.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 456  lr: 2.62144000000000017e-08  evaluation reward: 8
.94
episode: 3462  score: 8.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 439  lr: 2.62144000000000017e-08  evaluation reward: 8
.99
episode: 3463  score: 7.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 407  lr: 2.62144000000000017e-08  evaluation reward: 8
.98
episode: 3464  score: 9.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 421  lr: 2.62144000000000017e-08  evaluation reward: 9
.0
episode: 3465  score: 5.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 303  lr: 2.62144000000000017e-08  evaluation reward: 8
.97
episode: 3466  score: 8.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 407  lr: 2.62144000000000017e-08  evaluation reward: 8
.94
episode: 3467  score: 6.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 354  lr: 2.62144000000000017e-08  evaluation reward: 8
.95
episode: 3468  score: 4.0  memory length: 1000000  epsilon: 0.00999802000
8555413  steps: 279  lr: 2.62144000000000017e-08  evaluation reward: 8
.93
episode: 3469  score: 11.0  memory length: 1000000  epsilon: 0.0099980200
08555413  steps: 572  lr: 2.62144000000000017e-08  evaluation reward:
8.89
episode: 3470  score: 9.0  memory length: 1000000  epsilon: 0.00999802000
```



```
8555413      steps: 502      lr: 2.62144000000000017e-08      evaluation reward: 8
.9
episode: 3471      score: 12.0      memory length: 1000000      epsilon: 0.0099980200
08555413      steps: 483      lr: 2.62144000000000017e-08      evaluation reward:
8.93
episode: 3472      score: 5.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 283      lr: 2.62144000000000017e-08      evaluation reward: 8
.84
episode: 3473      score: 6.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 338      lr: 2.62144000000000017e-08      evaluation reward: 8
.84
episode: 3474      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 391      lr: 2.62144000000000017e-08      evaluation reward: 8
.88
episode: 3475      score: 5.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 289      lr: 2.62144000000000017e-08      evaluation reward: 8
.85
episode: 3476      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 440      lr: 2.62144000000000017e-08      evaluation reward: 8
.79
episode: 3477      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 404      lr: 2.62144000000000017e-08      evaluation reward: 8
.84
episode: 3478      score: 5.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 290      lr: 2.62144000000000017e-08      evaluation reward: 8
.83
episode: 3479      score: 2.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 198      lr: 2.62144000000000017e-08      evaluation reward: 8
.74
episode: 3480      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 375      lr: 2.62144000000000017e-08      evaluation reward: 8
.75
episode: 3481      score: 6.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 371      lr: 2.62144000000000017e-08      evaluation reward: 8
.75
episode: 3482      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 407      lr: 2.62144000000000017e-08      evaluation reward: 8
.79
episode: 3483      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 406      lr: 2.62144000000000017e-08      evaluation reward: 8
.8
episode: 3484      score: 7.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 387      lr: 2.62144000000000017e-08      evaluation reward: 8
.81
episode: 3485      score: 8.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 426      lr: 2.62144000000000017e-08      evaluation reward: 8
.75
episode: 3486      score: 3.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 213      lr: 2.62144000000000017e-08      evaluation reward: 8
.64
episode: 3487      score: 7.0      memory length: 1000000      epsilon: 0.00999802000
8555413      steps: 367      lr: 2.62144000000000017e-08      evaluation reward: 8
.62
```

episode: 3488	score: 7.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 366	lr: 2.62144000000000017e-08	evaluation reward: 8
.63			
episode: 3489	score: 3.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 226	lr: 2.62144000000000017e-08	evaluation reward: 8
.59			
episode: 3490	score: 7.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 405	lr: 2.62144000000000017e-08	evaluation reward: 8
.6			
episode: 3491	score: 11.0	memory length: 1000000	epsilon: 0.00999802000
08555413	steps: 540	lr: 2.62144000000000017e-08	evaluation reward:
8.63			
episode: 3492	score: 8.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 492	lr: 2.62144000000000017e-08	evaluation reward: 8
.64			
episode: 3493	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 484	lr: 2.62144000000000017e-08	evaluation reward: 8
.55			
episode: 3494	score: 13.0	memory length: 1000000	epsilon: 0.00999802000
08555413	steps: 584	lr: 2.62144000000000017e-08	evaluation reward:
8.57			
episode: 3495	score: 9.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 418	lr: 2.62144000000000017e-08	evaluation reward: 8
.59			
episode: 3496	score: 10.0	memory length: 1000000	epsilon: 0.00999802000
08555413	steps: 513	lr: 2.62144000000000017e-08	evaluation reward:
8.62			
episode: 3497	score: 6.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 324	lr: 2.62144000000000017e-08	evaluation reward: 8
.64			
episode: 3498	score: 7.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 383	lr: 2.62144000000000017e-08	evaluation reward: 8
.55			
episode: 3499	score: 8.0	memory length: 1000000	epsilon: 0.00999802000
8555413	steps: 421	lr: 2.62144000000000017e-08	evaluation reward: 8
.53			



## Visualize Agent Performance

BE AWARE THIS CODE BELOW MAY CRASH THE KERNEL IF YOU RUN THE SAME CELL TWICE.

Please save your model before running this portion of the code.

```
In [10]: torch.save(agent.policy_net, "./save_model/breakout_dqn_latest.pth")
```

```
In [ ]: from gym.wrappers import Monitor
        #from gym.wrappers.monitor import Monitor
        import glob
        import io
        import base64

        from IPython.display import HTML
        from IPython import display as ipythondisplay

        from pyvirtualdisplay import Display

        # Displaying the game live
        def show_state(env, step=0, info=""):
            plt.figure(3)
            plt.clf()
            plt.imshow(env.render(mode='rgb_array'))
            plt.title("%s | Step: %d %s" % ("Agent Playing", step, info))
            plt.axis('off')

            ipythondisplay.clear_output(wait=True)
            ipythondisplay.display(plt.gcf())

        # Recording the game and replaying the game afterwards
        def show_video():
            mp4list = glob.glob('video/*.mp4')
            if len(mp4list) > 0:
                mp4 = mp4list[0]
                video = io.open(mp4, 'r+b').read()
                encoded = base64.b64encode(video)
                ipythondisplay.display(HTML(data='''<video alt="test" autoplay
                    loop controls style="height: 400px;">
                    <source src="data:video/mp4;base64,{0}" type="video/mp4" />
                </video>'''.format(encoded.decode('ascii'))))
            else:
                print("Could not find video")

        def wrap_env(env):
            env = Monitor(env, './video', force=True)
            return env
```

```
In [ ]: display = Display(visible=0, size=(300, 200))
        display.start()

        # Load agent
        # agent.load_policy_net("./save_model/breakout_dqn.pth")
        agent.epsilon = 0.0 # Set agent to only exploit the best action

        env = gym.make('BreakoutDeterministic-v5')
        env = wrap_env(env)

        done = False
```

```

score = 0
step = 0
state = env.reset()
next_state = state
life = number_lives
history = np.zeros([5, 84, 84], dtype=np.uint8)
get_init_state(history, state)

while not done:

    # Render breakout
    env.render()
    # show_state(env, step) # uncommenting this provides another way to visual

    step += 1
    frame += 1

    # Perform a fire action if ball is no longer on screen
    if step > 1 and len(np.unique(next_state[:189] == state[:189])) < 2:
        action = 0
    else:
        action = agent.get_action(np.float32(history[:4, :, :]) / 255.)
        state = next_state

    next_state, reward, done, info = env.step(action + 1)

    frame_next_state = get_frame(next_state)
    history[4, :, :] = frame_next_state
    terminal_state = check_live(life, info['lives'])

    life = info['lives']
    r = np.clip(reward, -1, 1)
    r = reward

    # Store the transition in memory
    agent.memory.push(deepcopy(frame_next_state), action, r, terminal_state)
    # Start training after random sample generation
    score += reward

    history[:4, :, :] = history[1:, :, :]
env.close()
show_video()
display.stop()

```

In [ ]: